

In [1]: `import numpy as np
import pandas as pd
import seaborn as sns
import shap
df = pd.read_csv("HR.csv")
df.head()`

Out[1]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber	...	RelationshipSatisfaction	StandardHours	StockOptionLevel	TotalWorkingYear
0	41	Yes	Travel_Rarely	1102	Sales		1	2	Life Sciences	1	1	...	1	80	0
1	49	No	Travel_Frequently	279	Research & Development		8	1	Life Sciences	1	2	...	4	80	1
2	37	Yes	Travel_Rarely	1373	Research & Development		2	2	Other	1	4	...	2	80	0
3	33	No	Travel_Frequently	1392	Research & Development		3	4	Life Sciences	1	5	...	3	80	0
4	27	No	Travel_Rarely	591	Research & Development		2	1	Medical	1	7	...	4	80	1

5 rows × 35 columns

In [2]: `df.isnull().sum()
df["Attrition"].value_counts()`

Out[2]:

```
Attrition
No      1233
Yes      237
Name: count, dtype: int64
```

In [3]: `df.dropna(inplace=True)`

In [4]: `df.isnull().sum()`

Out[4]:

```
Age                0
Attrition          0
BusinessTravel     0
DailyRate         0
Department        0
DistanceFromHome  0
Education         0
EducationField     0
EmployeeCount     0
EmployeeNumber    0
EnvironmentSatisfaction 0
Gender            0
HourlyRate        0
JobInvolvement    0
JobLevel          0
JobRole           0
JobSatisfaction    0
MaritalStatus     0
MonthlyIncome     0
MonthlyRate       0
NumCompaniesWorked 0
Over18            0
OverTime          0
PercentSalaryHike  0
PerformanceRating  0
RelationshipSatisfaction 0
StandardHours     0
StockOptionLevel  0
TotalWorkingYears 0
TrainingTimesLastYear 0
WorkLifeBalance   0
YearsAtCompany    0
YearsInCurrentRole 0
YearsSinceLastPromotion 0
YearsWithCurrManager 0
dtype: int64
```

In [5]: `df.groupby("Department")["Attrition"].value_counts(normalize=True).unstack()`

Out[5]:

	Attrition	No	Yes
Department			
Human Resources		0.809524	0.190476
Research & Development		0.861602	0.138398
Sales		0.793722	0.206278

In [6]: `sns.boxplot(x="Attrition", y="MonthlyIncome", data=df)`

Out[6]:

<Axes: xlabel='Attrition', ylabel='MonthlyIncome'>

In [7]: `sns.histplot(data=df, x="YearsSinceLastPromotion", hue="Attrition", multiple="stack")`

Out[7]:

<Axes: xlabel='YearsSinceLastPromotion', ylabel='Count'>

In [8]: `from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score

Separate features and target
X = df.drop("Attrition", axis=1)
y = df["Attrition"].map({"Yes": 1, "No": 0}) # convert target to 0/1

One-hot encode categorical features
X = pd.get_dummies(X, drop_first=True)

Train-test split
X_train, X_test, y_train, y_test = train_test_split(
 X, y, test_size=0.2, random_state=42
)

Train Logistic Regression
lr = LogisticRegression(max_iter=5000, solver="lbfgs")
lr.fit(X_train, y_train)

Predictions
y_pred = lr.predict(X_test)

print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))
accuracy = accuracy_score(y_test, y_pred)
print("Model Accuracy:", accuracy)`

Confusion Matrix:

```
[[247  8]
 [ 23 16]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.91	0.97	0.94	255
1	0.67	0.41	0.51	39
accuracy			0.89	294
macro avg	0.79	0.69	0.72	294
weighted avg	0.88	0.89	0.88	294

Model Accuracy: 0.8945578231292517

C:\Users\ASUS\anaconda3\Lib\site-packages\sklearn\linear_model\logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(

In [9]: `# Ensure DataFrames
X_train_df = pd.DataFrame(X_train, columns=X.columns)
X_test_df = pd.DataFrame(X_test, columns=X.columns)

Use SHAP Explainer (works with linear models)
explainer = shap.Explainer(lr, X_train_df)

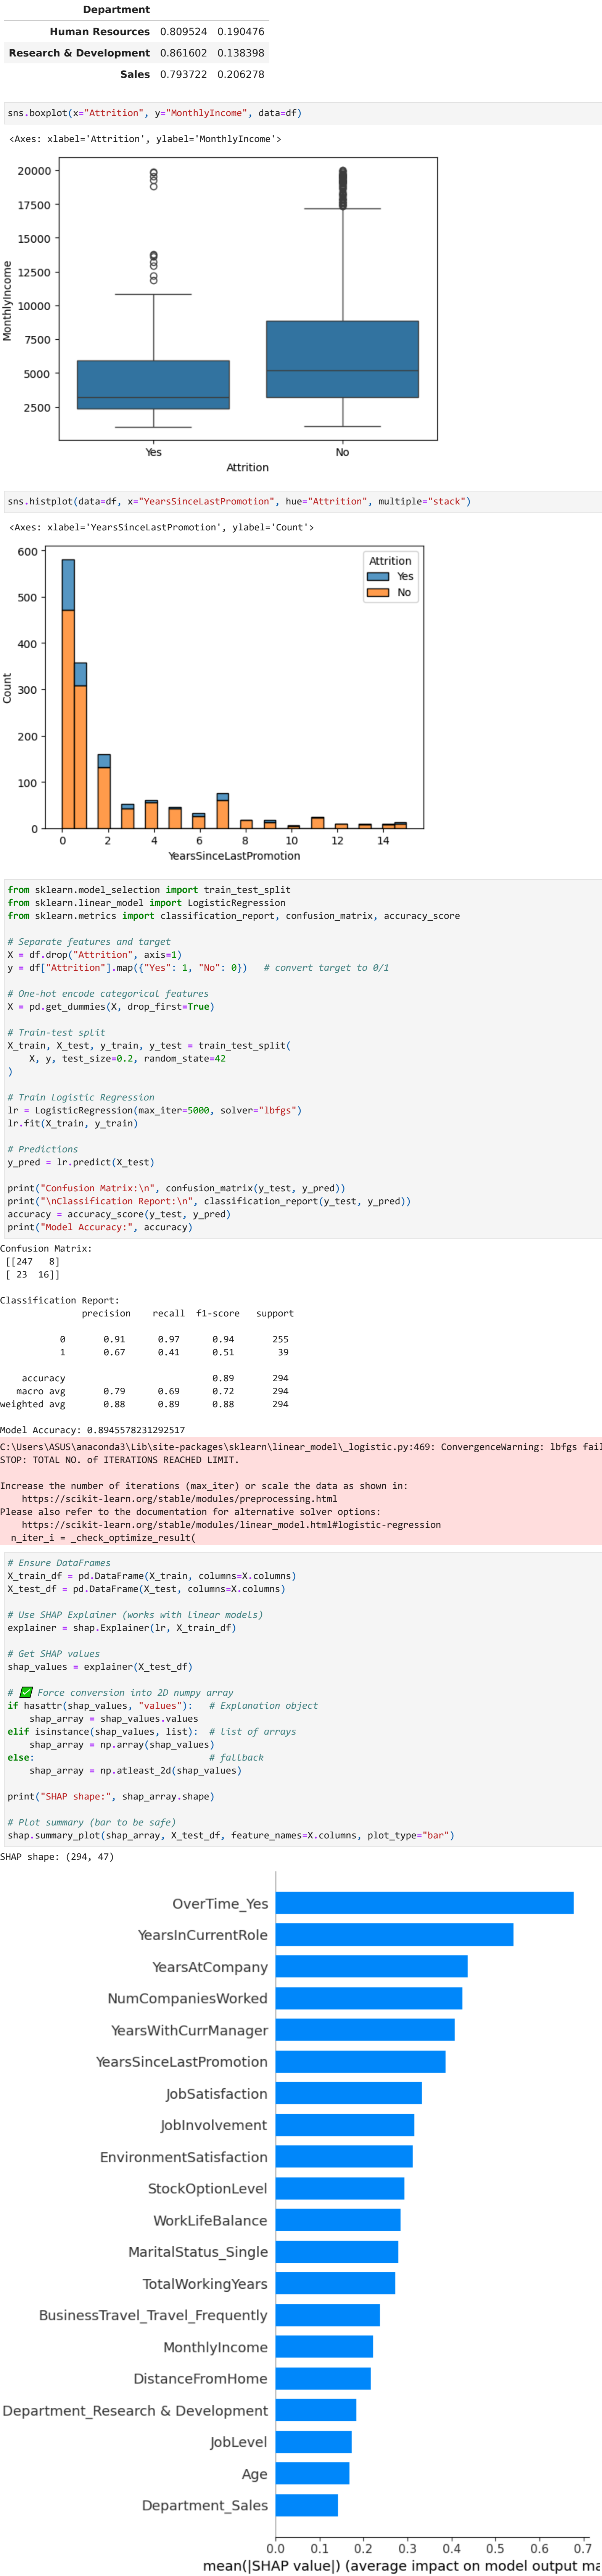
Get SHAP values
shap_values = explainer(X_test_df)

Force conversion into 2D numpy array
if hasattr(shap_values, "values"): # Explanation object
 shap_array = shap_values.values
elif isinstance(shap_values, list): # List of arrays
 shap_array = np.array(shap_values)
else: # fallback
 shap_array = np.atleast_2d(shap_values)

print("SHAP shape:", shap_array.shape)

Plot summary (bar to be safe)
shap.summary_plot(shap_array, X_test_df, feature_names=X.columns, plot_type="bar")`

SHAP shape: (294, 47)



```
In [10]: df["Attrition_Prob"] = lr.predict_proba(X[:,1])
df.to_csv("hr_attrition_results.csv", index=False)

In [ ]:
```